

Evolutionary optimization of machine learning pipelines

Gabriela Suchopárová

Abstract

We have developed a system that, for a given supervised learning task, finds a suitable pipeline — combination of machine learning, ensembles and preprocessing methods — using the developmental genetic programming.

AutoML

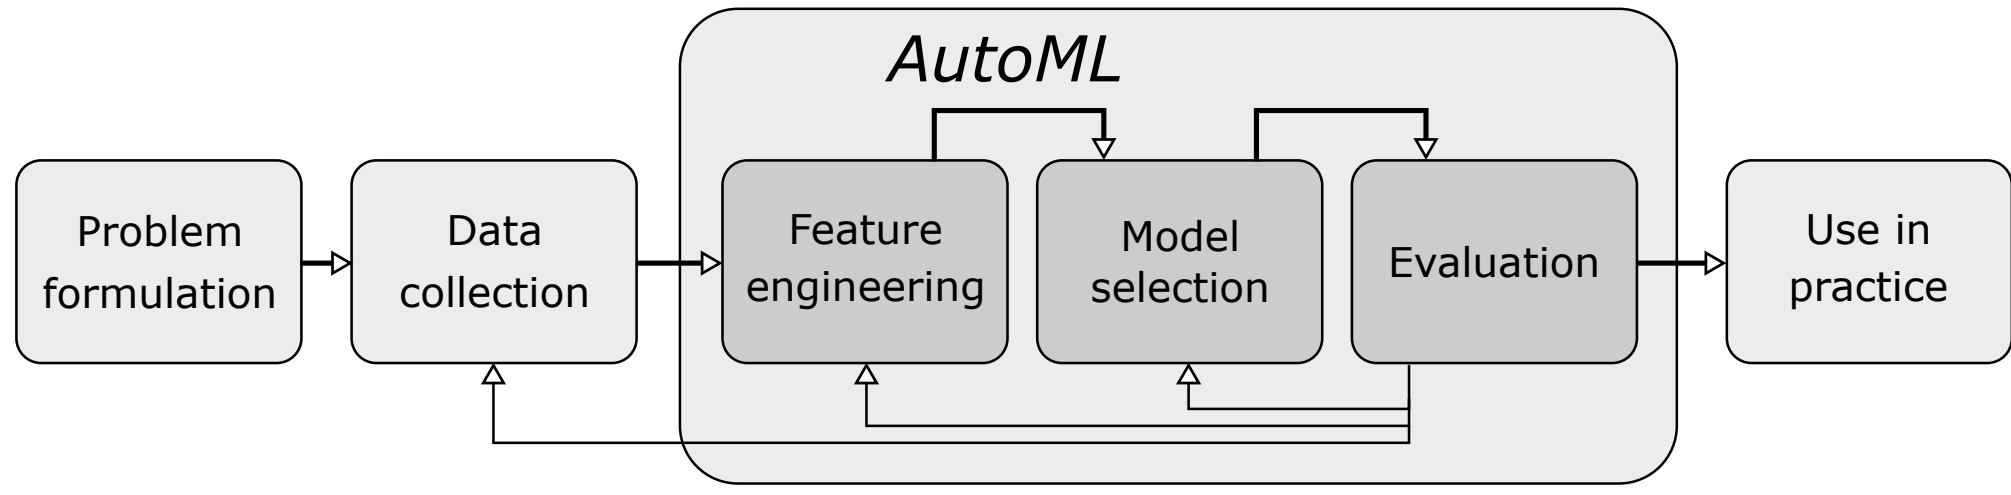


Figure 1: Schema of a general workflow

The AutoML systems aim to automatize a part of the machine learning workflow, which is the process of selection of suitable feature preprocessing methods and of a model with a particular hyperparameter setting.

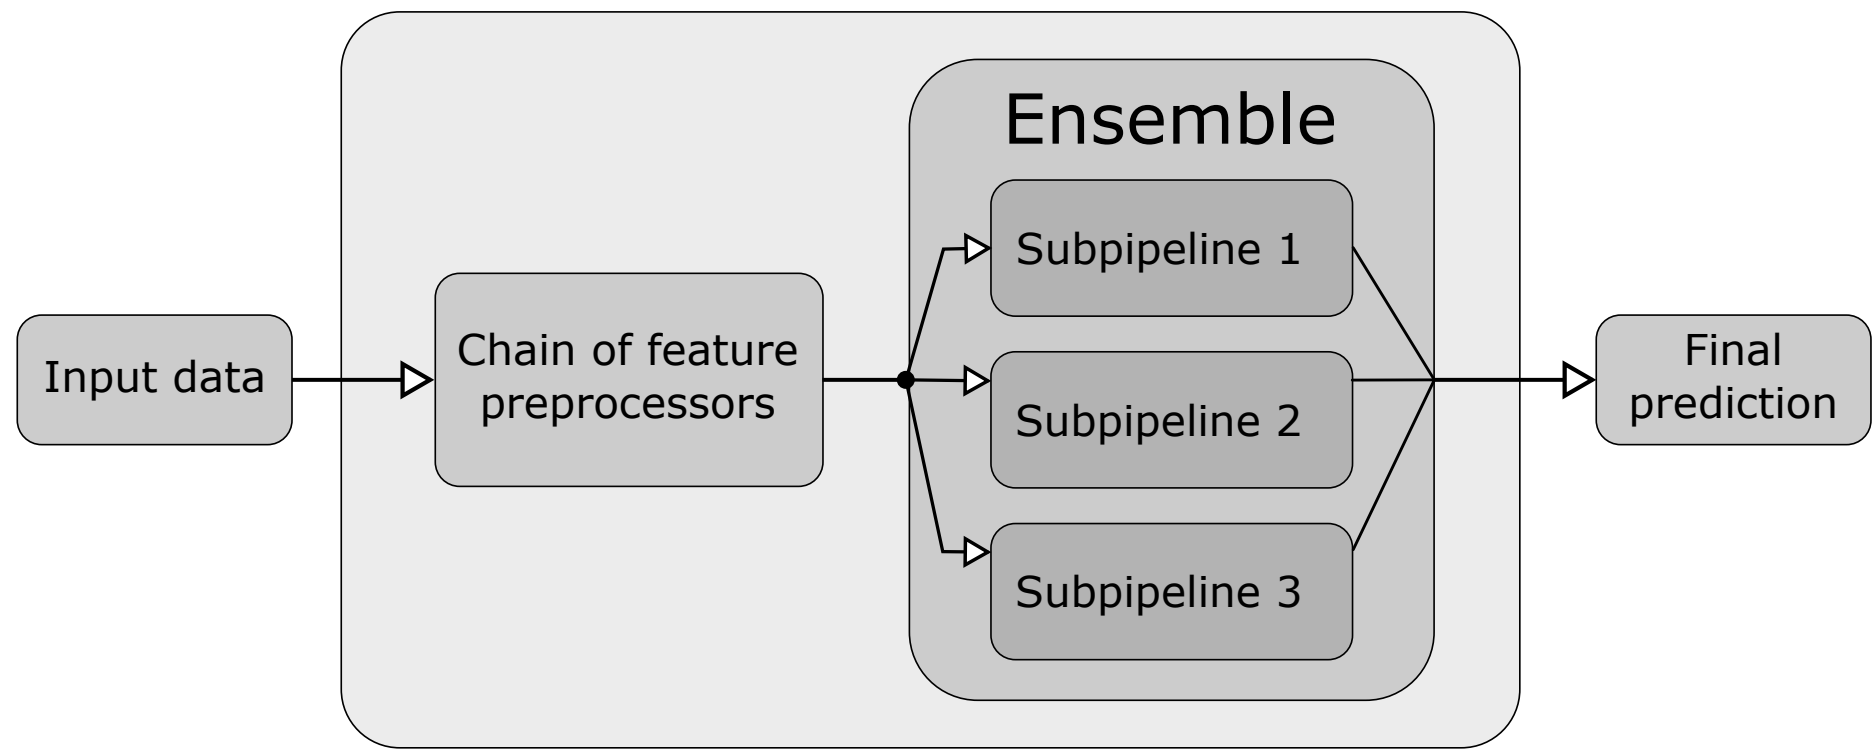


Figure 2: A more complex machine learning pipeline

One of the tasks targeted by AutoML systems is *pipeline optimization*. A machine learning pipeline is composed from a chain of feature preprocessing methods and one estimator. The estimator can be a simple method or an ensemble that contains one or more base-estimators. Because a base-estimator can be a pipeline as well, a general pipeline is a directed acyclic graph.

Our system

Our system aims to search for a pipeline that performs well on the given task. A pipeline can be composed from both ensembles and feature preprocessing methods, but simple methods are tried out as well.

Developmental GP

For pipeline optimization, we used the genetic programming (GP), which is a subfield of evolutionary algorithms.




Figure 3: Example of an encoded pipeline

We created a specific encoding that enables to convert pipelines in the form of a DAG into a tree representation. Instead of directly encoding pipeline steps as nodes, we apply the developmental GP, where the nodes represent *operations* that create the pipeline. An example of the encoding is shown in Figure 3. The tree individual contains instructions which construct the actual pipeline:

- cPipe** — create a pipeline with a preprocessor chain and an estimator
 - **ada** — insert an AdaBoost ensemble with a base-estimator
 - **cPipe** — create a pipeline with a preprocessor chain
 - **SGD** — insert a SGD classifier
 - **MinMaxScaler** — insert a MinMaxScaler
 - **PCA** — insert the PCA preprocessor

OpenML experiment

Our system was tested on the OpenML-CC18 benchmark, which is a suite of 72 datasets available at OpenML. Each of the datasets is associated with a classification task, the evaluation method is 10-fold cross-validation. We present a comparison of our results with all other results uploaded to OpenML. On most of the datasets our system performed in the upper quartile of all results.

OpenML-CC18 benchmark

